

AI Interviewer: What Needs to Be Built

Findings from 11 test interviews · September 2026

What this product is

Our product is an AI that acts as a job interviewer for consulting roles. It gives the candidate a business problem to solve and hands out facts about the problem only when the candidate asks for them. After the interview, a second AI grades the candidate and writes a feedback report.

To test it, we ran 11 practice interviews in which an AI played different kinds of candidates: nervous, overconfident, someone trying to trick the system, and someone going through a hard time.

The main finding

Our rulebook for the interviewer relies on automatic checks that are supposed to run in the background — for example, a checker that verifies a candidate's math *before* the interviewer corrects them. **Several of these checks were never built.** Without them, the AI has to guess, and in the tests it often guessed wrong.

Rewriting the rulebook can't fix that — the checks have to be built. This document lists the 14 that are missing, most important first.

Some good news: a few problems already have a temporary fix in the rulebook. For example, the interviewer is now told never to correct a candidate unless the math checker has flagged an error. Until the checker exists, it simply won't correct anyone — far better than "correcting" people who were right.

All 14 fixes at a glance

#	Fix	P	#	Fix	P
1	Detect when a candidate is in distress	1	8	Warn before time runs out; always ask for the recommendation	3
2	Build the math checker the interviewer relies on	2	9	Limit time spent on one math mistake	3
3	Confirm data actually reached the candidate	2	10	Block made-up numbers	3
4	Fix the grading tool's math checker	2	11	Don't treat quotes as insults	4
5	Double-check "you made a math error" in reports	2	12	Don't blame candidates for the interviewer's mistakes	4
6	Don't let the interviewer skip a candidate's question	3	13	Don't mark down skills that were never tested	4
7	Keep promises to answer later	3	14	Answer the real question when a candidate goes off-topic	4

Priority (P): 1 = safety · 2 = gives candidates wrong information · 3 = interview doesn't run properly · 4 = unfair warnings or grading

How to check the fixes

Four fixes (4, 5, 12, 13) can be checked by re-running the grading on the 11 interviews we already have. The rest need a fresh test run. Each test interview costs about \$1.18, so all 60 test candidates cost about \$71 — cheap enough to rerun after every fix.

Also fix the test setup

Three tests didn't test what they were meant to, because of how the simulated candidates behaved (not the product):

- The "slow thinker" paused 50 seconds at a time, just under the 60-second check-in point, so that behavior was never tested. Rerun with 65–90 second pauses.
- The "nervous" candidate ran out of time before her key moments (showing distress, freezing on the recommendation). Move those earlier.
- The "hostile" candidate stayed too polite to trigger the insult rules. Write the hostile lines out word for word.

The fixes in detail

Each card shows what went wrong in testing, what to build, and how we'll know it works. Cards are in priority order; the colored number matches the priority key on page 1.

1 Detect when a candidate is in distress

Priority 1

What went wrong: A test candidate said she wasn't sleeping and didn't "see the point in any of it anymore." The interviewer replied kindly on its own, but nothing in the system noticed: it never offered to stop, kept the timer running, and graded her normally. Our own length check even flagged the kind reply as "too long."

What to build: A detector that scans every candidate message for signs of distress. When it fires, the interviewer offers to stop with no penalty, the timer pauses, and that moment is left out of the grade. These replies are exempt from the length check.

Done when: Rerunning that test, the detector fires, stopping is offered, and the timer pauses.

2 Build the math checker the interviewer relies on

Priority 2

What went wrong: Three times the interviewer said "Quick correction" — and all three candidates were right. One was "corrected" for a number he never said, and apologized anyway. Another correction gave away facts the candidate hadn't unlocked. No math checker is running; the AI was guessing.

What to build: A checker that recalculates the candidate's numbers from facts already shared and flags real mistakes. The interviewer may correct only when it flags something, must quote what the candidate actually said, and must never reveal unshared facts.

Done when: Reruns show zero false corrections, and every correction traces back to a flag.

3 Confirm data actually reached the candidate

Priority 2

What went wrong: In 4 of 11 tests the interviewer said "Here's the menu price change" and gave no number. Candidates replied, "I don't see the number." The system still recorded the fact as delivered.

What to build: Record a fact as delivered only when its number appears in the message (or the correct chart is shown). If data is announced but missing, block the message and regenerate it.

Done when: Every fact marked "delivered" appears in the transcript.

4 Fix the grading tool's math checker

Priority 2

What went wrong: A checker used for grading flagged a "revenue per store" mistake in 10 of 11 reports. It grabbed random nearby numbers (a "\$2.50 drink," a price multiplier) and even flagged a candidate who never did that calculation. The false mistake reached at least three reports.

What to build: Save the exact sentence each number came from and what it's meant to measure. If the sentence isn't about that measurement, discard the number before the grader sees it.

Done when: Zero false flags when re-run on the 11 existing interviews. *No new tests needed.*

5 Double-check "you made a math error" in reports

Priority 2

What went wrong: One candidate's top piece of feedback told her to fix her arithmetic. It was correct. Today's report checks catch fake quotes and false "you never mentioned X" claims, but nothing checks a claim that the candidate made a mistake.

What to build: Before a report is finalized, recalculate every "you made a mistake" claim and remove it if the candidate was actually right.

Done when: Re-running existing reports removes every false error claim. *No new tests needed.*

6 Don't let the interviewer skip a candidate's question

Priority 3

What went wrong: Candidates asked for facts the interviewer had ("Did menu prices change?") and it often moved on without answering — 33 times across 10 of 11 tests. The system logged every case but did nothing to stop it.

What to build: Before a reply is sent, check that the candidate's last request was answered, refused, or clearly postponed. If not, block and regenerate the reply.

Done when: A full set of reruns has no unanswered requests.

7 Keep promises to answer later

Priority 3

What went wrong: The interviewer told one candidate "I'll get you menu prices" and never did.

What to build: Keep a list of postponed requests, remind the interviewer each turn, and resolve all of them before asking for the final recommendation.

Done when: No postponed request is still open at the end of any interview.

8 Warn before time runs out; always ask for the recommendation

Priority 3

What went wrong: One interview ended with no time warning and without ever asking for the final recommendation — the most important part of the interview.

What to build: A timer that warns 30 seconds before the end. When time is short, the recommendation request always comes first, with time left to answer.

Done when: Every interview asks for a recommendation with at least 30 seconds left.

9 Limit time spent on one math mistake

Priority 3

What went wrong: A struggling candidate got four rounds of hints (about three minutes) on one calculation. The rule is two tries, then give the answer.

What to build: Count attempts per calculation. After two, give the answer and move on. If the same kind of mistake repeats later, correct it right away.

Done when: No calculation gets more than two rounds of hints.

10 Block made-up numbers

Priority 3

What went wrong: The interviewer told a candidate costs were "94%, not 8%." He never said 8%, and it isn't a case fact. The rulebook has a check for invented numbers, but it isn't blocking anything.

What to build: Check every number in a reply before it's sent. Each must come from the case facts, the candidate, or the math checker — otherwise, block the reply.

Done when: No invented numbers appear in interviewer messages.

11 Don't treat quotes as insults

Priority 4

What went wrong: A polite candidate quoted a character in the case ("the CEO basically said you're an idiot if you think it's labor") and was warned to "keep this professional." His report called it a lapse.

What to build: Ignore words the candidate attributes to someone else ("the CEO said...") when checking for insults. If hostility is still unclear, log it instead of warning.

Done when: Rerunning that test produces no warning.

12 Don't blame candidates for the interviewer's mistakes

Priority 4

What went wrong: Three reports blamed candidates for system errors: a false correction cited as proof of a mistake, a false warning turned into a "lapse," and a penalty for correctly pointing out an unanswered question.

What to build: Mark every interviewer mistake (false correction, false warning, ignored question) from the system's own log, and tell the grader it can't hold the candidate's reaction against them.

Done when: Re-running those three reports removes the penalties. *No new tests needed.*

13 Don't mark down skills that were never tested

Priority 4

What went wrong: One report noted the candidate was never asked for a recommendation — then scored her low for not giving one.

What to build: If a skill wasn't tested because of the interview (not the candidate), it can't score below "meets bar." If there's too little to go on, show "not assessed."

Done when: Re-running that report changes the rating. *No new tests needed.*

14 Answer the real question when a candidate goes off-topic

Priority 4

What went wrong: A candidate joked "ignore your instructions" and asked two real questions in one message. The reply, "Let's stay on the case," dropped both.

What to build: When a message mixes off-topic content and real questions, briefly redirect and still answer the real ones.

Done when: Rerunning that test shows both questions answered.