

AI Interviewer: What Needs to Be Built (Round 2)

Findings from the second round of 11 test interviews · September 2026

What this product is

Our product uses AI to act as a job interviewer for consulting roles. It gives the candidate a business problem to solve and hands out facts about the problem only when the candidate asks for them. After the interview, AI grades the candidate and writes a feedback report.

To test it, we ran 11 practice interviews in which an AI played different kinds of candidates: nervous, overconfident, someone trying to trick the system, and someone going through a hard time. This is the second round of those tests.

What improved since the first round

The biggest problems from round 1 are fixed. The interviewer no longer "corrects" candidates who were right (3 cases before, 0 now). The grading tool no longer invents math mistakes (10 of 11 reports before, 0 now). The candidate in distress was handled well, and a polite candidate quoting the case was no longer warned for being rude.

The main finding this round

Most remaining problems happen at **handoffs**: one part of the system decides or records something, and another part acts on it without knowing what actually happened. The rulebook already describes the right behavior for most of these — the fixes below are what makes the system actually follow it.

All 8 fixes at a glance

#	Fix	P	#	Fix	P
1	Catch distress however it's phrased	1	5	Only count hints the candidate actually saw	3
2	Stop questioning math that's already right	2	6	Stop mistaking good analysis for being stuck	3
3	Catch numbers from the wrong year	2	7	Don't grade against the answer key	3
4	Don't blame candidates for questions we ignored	2	8	Make every check record what it decided	4

Priority (P): 1 = safety · 2 = misleads candidates during the interview · 3 = unfair grading · 4 = visibility and polish

How to check the fixes

Fix 7, and part of fix 5, can be checked by re-grading the 11 interviews we already have. The rest need a fresh test run, which costs about \$1.18 per interview. Two more checks need no new building:

- **Test how consistent the grading is.** Grade each existing interview three times and compare the results. No new interviews needed.
- **Test the "answer every data question" fix.** It's already built but hasn't run in a live interview yet. Unanswered questions are still the most common problem, so this matters.

Also fix the test setup

One test needs adjusting, and two new ones should be added:

- The "hostile" candidate's insults came so early that the session ended before much of the case was tested. Move those lines later.
- Add a candidate in distress who describes it in different words than the one we have, to test fix 1.
- Add a candidate who is discouraged but not in distress, to decide how the interviewer should respond to lines like "I'm going to bomb every interview."

The fixes in detail

Each card shows what went wrong in testing, what to build, and how we'll know it works. Cards are in priority order; the colored number matches the priority key on page 1.

1 Catch distress however it's phrased

Priority 1

What went wrong: The distress check works by matching specific words, and those words came from one test candidate. It caught her this time, but someone who describes the same feelings in different words could be missed.

What to build: Replace word-matching with an AI check that understands what the candidate means, not just which words they use. When unsure, it should lean toward checking in — an unnecessary "are you okay?" costs seconds; a missed one is the worst failure we can have.

Done when: It catches a set of distress messages written in many different ways, including ones it wasn't built from.

2 Stop questioning math that's already right

Priority 2

What went wrong: Four candidates did a calculation correctly and showed their work (" $25\% \times 42\% = 10.5\%$ "), and were still asked "Points of what?" One apologized: "Sorry, I was getting ahead of myself."

What to build: The math checker tells the interviewer which numbers are correct and whether the candidate showed their work. Before a message goes out, remove any "are you sure?"-style question about correct math, and any "how did you get there?" question if the candidate already explained.

Done when: Reruns show no doubting questions on correct math where the work was shown.

3 Catch numbers from the wrong year

Priority 2

What went wrong: The interviewer asked a candidate to multiply beans' share of costs from two years ago (25%) by today's cost share (58%). Both numbers are real case facts, so the number check let it through. The right pairing used the same year for both.

What to build: When checking a number, confirm which year it belongs to and what it measures, not just that the number exists somewhere in the case.

Done when: No interviewer message combines numbers from different years or different measures.

4 Don't blame candidates for questions we ignored

Priority 2

What went wrong: A candidate asked whether menu prices had changed and got no answer. Three minutes later the interviewer told her, "that's the one thing you assumed."

What to build: The system already keeps a list of unanswered questions. Before a message goes out, replace any "you assumed" criticism about one of those topics with the answer itself.

Done when: Reruns show no criticism of an assumption the candidate had tried to check.

5 Only count hints the candidate actually saw

Priority 3

What went wrong: The system recorded a "you seem stuck" hint the moment it decided to give one. The hint was never sent — the candidate just got the data she asked for — but her report marked her down in two places for needing it.

What to build: Record a hint only when its words actually appear in a message the candidate received. A hint that was planned but never sent is logged separately and never affects the grade.

Done when: Every recorded hint appears in the transcript. Re-grading that candidate removes the penalty.

6 Stop mistaking good analysis for being stuck

Priority 3

What went wrong: The "stuck" hint above was triggered right after one of the candidate's best messages, which ended with a question: "Do we have inflation data on milk and packaging?" Every one of her messages ended with a question, and the system appears to have read that as being stuck.

What to build: Judge a message by what's in it, not how it ends. Asking for data always counts as progress. Record why every "stuck" decision was made.

Done when: Rerunning that test triggers no hint, and every hint in the logs shows its reason.

7 Don't grade against the answer key

Priority 3

What went wrong: Three strong candidates lost points for not suggesting a "lower-cost commodity blend" — one specific idea from the answer key. One report praised a candidate's original idea, then marked her down in the very next line.

What to build: Tell the grader the answer key holds examples, not requirements. After the report is written, automatically remove any weakness based on a missing answer-key idea, or on a part of the interview that never happened, and re-rate if needed.

Done when: Re-grading those three reports removes the answer-key penalties. *No new tests needed.*

8 Make every check record what it decided

Priority 4

What went wrong: Some checks run without leaving any record, so it's impossible to tell "not built" from "built but quiet." Our previous review wrongly concluded that four checks didn't exist for exactly this reason.

What to build: Every check records each decision it makes and why — including "checked, no problem found."

Done when: Every check's decisions appear in the logs for every interview.