

AI Interviewer: What Needs to Be Built (Round 3)

Findings from the third round of 10 test interviews · October 2026

What this product is

Our product uses AI to act as a job interviewer for consulting roles. It gives the candidate a business problem to solve and hands out facts about the problem only when the candidate asks for them. After the interview, AI grades the candidate and writes a feedback report.

To test it, we ran 10 practice interviews in which an AI played different kinds of candidates. This round reran the hardest cases from round 2 — the nervous candidate, the hostile one, the one in distress, the non-native speaker — and added two new ones: a candidate in distress who describes it in new words, and a candidate who is discouraged but fine.

What improved since the second round

The round-2 fixes held up in live interviews. The nervous candidate heard **one goodbye** instead of four. Questions about facts the interviewer had went unanswered **4 times instead of 23**. Candidates who did math correctly were **never questioned on it** (5 times before). The false "you seem stuck" hint is gone. The new distress check caught a candidate whose words the old check would have missed, and the hostile candidate was only ended after most of the case had been tested.

The main finding this round

What's left is mostly about **proportion**: the interviewer overreacted to a joke about interviews, and the grader rates almost everyone "strong." Two problems turned up with the candidate who never gives a recommendation; both were fixed right after the round and need one more live check.

All 6 fixes at a glance

#	Fix	P	#	Fix	P
1	Don't treat a joke about interviews as a crisis	1	4	Answer a question asked in the last message	4
2	Deliver every fact the interviewer announces	2	5	Keep system words out of the interviewer's mouth	4
3	Make the grades spread out	3	6	Confirm that recorded hints were really hints	4

Priority (P): 1 = safety · 2 = misleads candidates during the interview · 3 = unfair grading · 4 = visibility and polish · A, B = already fixed, needs a live check

How to check the fixes

Fix 3 can be checked on interviews we already have. The rest need a fresh test run, which cost about **\$1.03 per interview** this round (\$10.33 for all 10). Two checks need no new building:

- **Test how consistent the grading is.** Grade each existing interview three times and compare. Still not done, and it matters more now that fix 3 changes the grading.
- **Rerun the nervous candidate.** Fixes A and B were made after this round and have only been checked against her old transcript.

Also fix the test setup

- Add a deliberately **average** candidate, so we can see whether fix 3 works.
- Add the "context confuser" candidate, who uses real case numbers in the wrong places. Round 2's "numbers from the wrong year" check is built, but had nothing to catch: on review, the round-2 example used the candidate's own guess, not a case fact.

The fixes in detail

Each card shows what went wrong in testing, what to build, and how we'll know it works. Cards are in priority order; the colored number matches the priority key on page 1.

1 Don't treat a joke about interviews as a crisis

Priority 1

What went wrong: A candidate frustrated only with this case said lightly, "honestly I'm going to bomb every interview at this rate." The interviewer stopped the case, asked if he was okay, and gave him a crisis hotline number. He had to reassure it: "I didn't mean to worry you... just me grumbling at myself about this case." The old word-matching check flags that phrase on purpose and runs first, so the new AI check — which reads meaning, and would have called it ordinary frustration — never got a say.

What to build: Take that phrase off the word list and let the AI check decide. Keep the clear phrases (self-harm, "I don't see the point in any of it anymore") on the list. Give the crisis line only when something suggests risk to the person, not for frustration about interviews.

Decision needed: this trades a small risk of missing real distress for fewer false alarms. The AI check caught 16 of 16 distress messages in testing, including one in this round that the word list missed.

Done when: The discouraged test candidate gets no check-in, and the distressed candidates still do.

2 Deliver every fact the interviewer announces

Priority 2

What went wrong: The interviewer said "here's the bean price change and the other-input change," then gave only the bean figure. The other figure arrived three minutes later. The safety net for announced-but-missing facts only runs when no fact was delivered at all.

What to build: Check each announced fact on its own, and add any missing one to the same message.

Done when: Every fact the interviewer announces appears in that same message.

3 Make the grades spread out

Priority 3

What went wrong: 8 of 9 graded interviews were rated "strong" overall. One candidate was written to do only adequate analysis and got "strong" on all eight skills. When almost everyone is strong, the report can't show candidates what to fix, and we can't test whether grading is fair across candidate types.

What to build: Add concrete examples to the grading rules of what separates "strong" from "meets the bar" on each skill.

Done when: Average test candidates mostly land at "meets the bar," and grading the same interview three times gives the same result.

4 Answer a question asked in the last message

Priority 4

What went wrong: A candidate asked for store-count figures in her final message, and the interview closed without answering. Her grade correctly didn't count it against her, but she never got an answer.

What to build: When the interview ends, answer or decline any request in the candidate's last message before the goodbye.

Done when: No request in a final message goes unanswered.

5 Keep system words out of the interviewer's mouth

Priority 4

What went wrong: A candidate asked whether the client had made big strategic changes. The interviewer answered, "I don't have anything flagged on that." "Flagged" is the system's own language and sounds like the interviewer is reading from a screen.

What to build: Add system words like this to the filter that already removes the interviewer's internal planning from its messages, and remind the interviewer to say "that's not in the information I have."

Done when: No system terms appear in transcripts.

6 Confirm that recorded hints were really hints

Priority 4

What went wrong: After fix B below, any question the interviewer asks at a moment when a hint is due is recorded as the hint. That's almost always right, but if the interviewer asked something unrelated instead, the candidate would be marked as having needed help.

What to build: A small AI check that confirms the message actually contains help, run only when a hint is due.

Done when: Every recorded hint, read by a person, is actually a hint.

Already fixed after this round — needs a live check

Both problems came from the nervous test candidate, who is written to answer "I don't know" every time she's asked for a recommendation. Both were fixed the same day and checked against her transcript, not yet in a live interview.

A Stop asking for the recommendation over and over

Fixed

What went wrong: Each time the interviewer tried to wrap up, the system blocked the ending because no recommendation had been given, and asked for one again. She was asked 17 times in about three minutes, until the clock ran out.

What we changed: After two refused asks, the interview may end without a recommendation, recorded as the candidate's result — which is what the rulebook already says.

Check in the next run: Her interview ends soon after the second refusal.

B Count hints the interviewer words in its own way

Fixed

What went wrong: The round-2 rule "only count hints the candidate actually saw" looked for set phrases. The interviewer's real hints — "Stay with it. Of those three, which do you want to see first?" — weren't recognized, so the system treated them as never sent, and her hints never moved past the first level.

What we changed: A hint counts if the message asks the candidate something. A message with no question (just data) still doesn't count.

Check in the next run: Her hints step up from level 1 to level 2 when she stays stuck.